



Software Track - SPUMA benchmarking on 236M cells occDrivAerStaticMesh test case

01 / 07 / 2025

Simone Bnà (s.bn@cineca.it)

Giuseppe Giaquinto (g.giaquinto@cineca.it)

Tommaso Zanelli (t.zanelli@cineca.it)

- **SPUMA** (**S**imulation **P**rocessing on **U**nified **M**emory **A**ccelerators) is a **minimally invasive** porting of OpenFOAM to GPU under development by Cineca.
- It is **portable** across different hardware (e.g. x86, Nvidia, AMD)
- Native OpenFOAM smoothers (Gauss-Seidel, DIC) cannot be trivially parallelized on GPU accelerators
- Additional algorithms were implemented to provide **scalable** parallel smoothers
 - Chebyshev
 - Two-stage Gauss-Seidel
 - Richardson

Chebyshev polynomial smoother

Given a polynomial degree n , the $n + 1$ approximate solution to $\mathbf{A} \mathbf{x} = \mathbf{b}$ is given by:

$$\mathbf{x}^{n+1} = \mathbf{x}^n + \alpha_0^n (\mathbf{x}^n - \mathbf{x}^{n-1}) + \alpha_1^n \mathbf{D}^{-1} (\mathbf{b} - \mathbf{A} \mathbf{x}^n)$$

$$\text{Where: } \mathbf{D} = \text{diag}(\mathbf{A}), \alpha_0^0 = 0, \alpha_1^0 = \rho^0, \alpha_1^n = \frac{4\rho^n}{\lambda_{\max} - \lambda_{\min}},$$

$$\rho^0 = \frac{\lambda_{\max} - \lambda_{\min}}{\lambda_{\max} + \lambda_{\min}} \text{ and } \rho^n = \left(2 \frac{\lambda_{\max} + \lambda_{\min}}{\lambda_{\max} - \lambda_{\min}} - \rho^{n-1} \right)^{-1}$$

The **Chebyshev polynomial** is constructed to strongly damp eigenvalues in the range $[\lambda_{\min}, \lambda_{\max}]$.

In the actual implementation, the value of λ_{\max} is either set a priori or estimated via the **Gershgorin theorem**, while λ_{\min} is estimated as

$$\lambda_{\min} = \frac{1}{8} \lambda_{\max}$$

Two-stage Gauss-Seidel smoother

In the **Gauss-Seidel** method the approximate solution to the linear system $\mathbf{A} \mathbf{x} = \mathbf{b}$ is given, at each iteration, by:

$$\mathbf{x}^{n+1} = \mathbf{x}^n + \mathbf{L}^{-1} (\mathbf{b} - \mathbf{A} \mathbf{x}^n)$$

Where \mathbf{L} is the lower triangular portion of \mathbf{A} .

\mathbf{L}^{-1} can be approximated as:

$$\sum_{j=0}^s (-\mathbf{D}^{-1} \mathbf{L})^j \mathbf{D}^{-1}$$

For $s = 1$ the **two-stage Gauss-Seidel** method is obtained, and the $n + 1$ iteration of \mathbf{x} is given by:

$$\mathbf{x}^{n+1} = \mathbf{x}^n + (\mathbf{I} - \mathbf{D}^{-1} \mathbf{L}) \mathbf{D}^{-1} (\mathbf{b} - \mathbf{A} \mathbf{x}^n)$$

HARDWARE USED

CINECA

- Benchmarking tests were run on Cineca's Leonardo cluster.
- Leonardo partitions are:

Booster

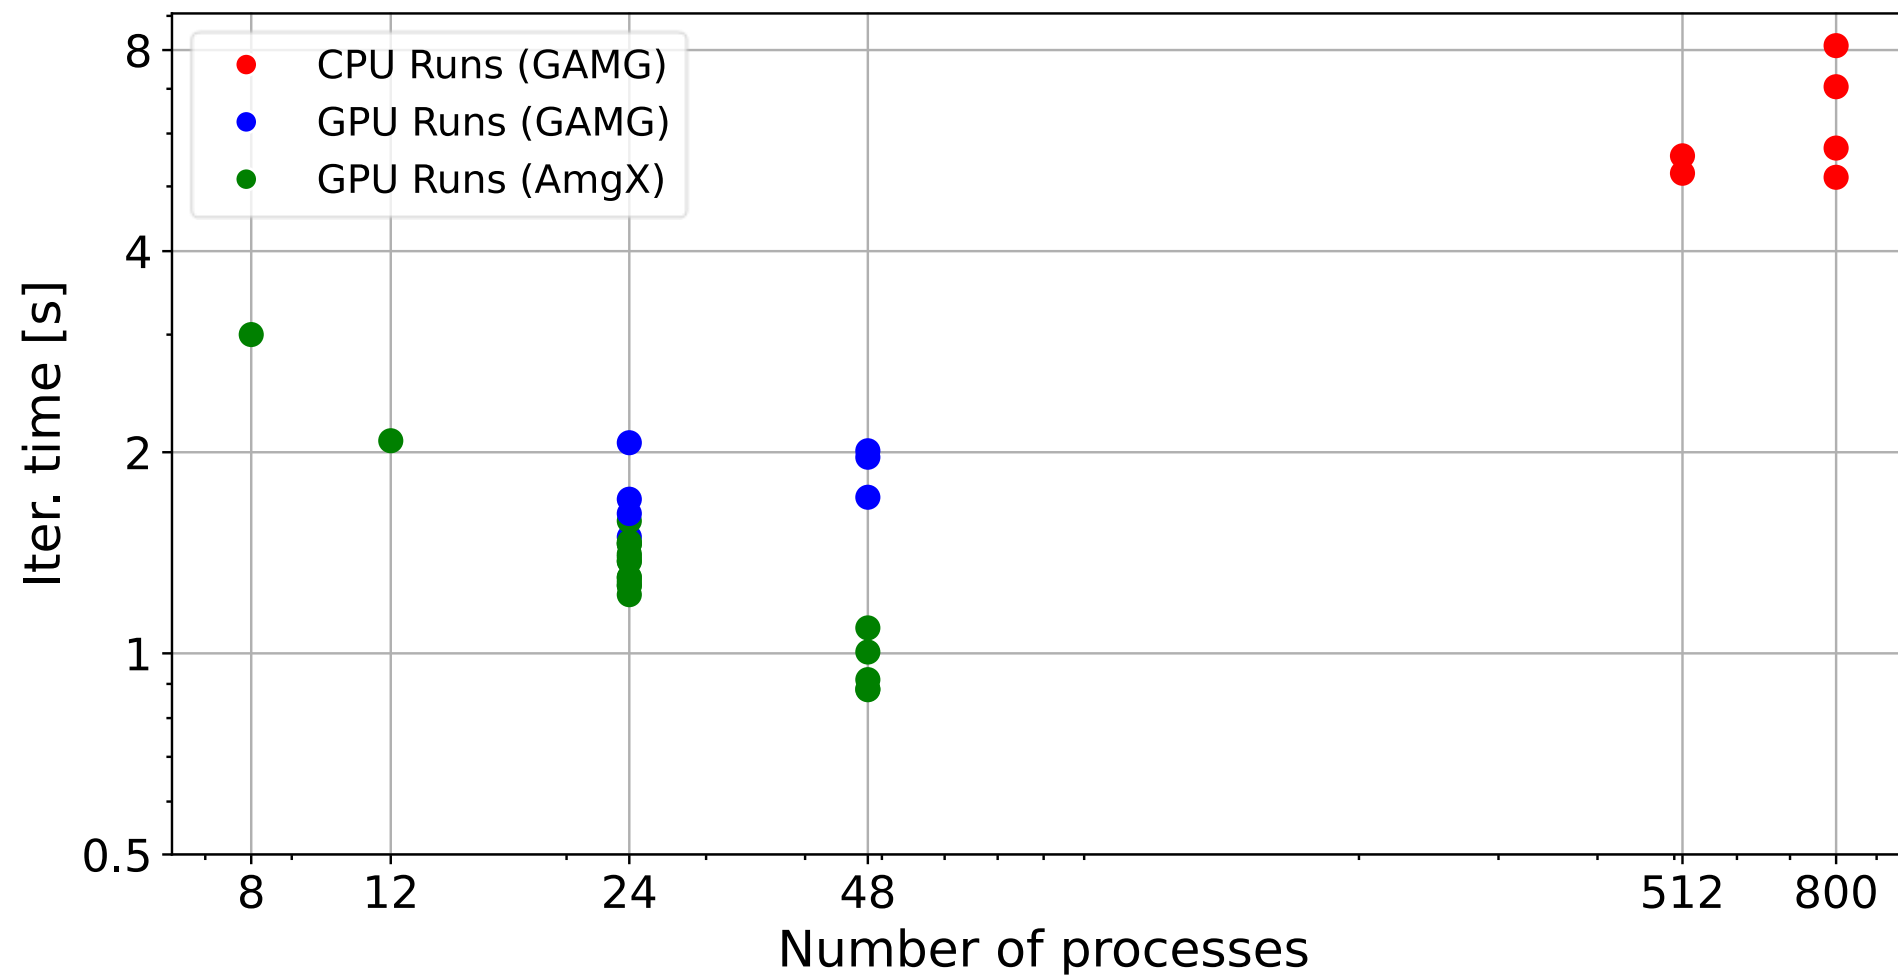
- **116** BullSequana XH2000 Direct Liquid cooling racks
 - **240** PFLOPs HPL Linpack Performance (Rmax)
- **3456** computing node servers
 - **13824 NVIDIA Ampere GPUs (A100)**
 - **884** TB of HBM2

Data Centric

- 16** BullSequana XH2000 Direct Liquid cooling racks
 - 9** PFLOPs HPL Linpack Performance (Rmax)
- 1536** computing node servers
- 3072** Intel **Sapphire Rapids 8480+** CPUs
- 786** TB DDR5 Memory
- 5.8** PB of local NVM

BENCHMARK RESULTS

Summary



Notes

- Fixed iteration runs took consistently longer than **fixed tolerance** runs.
- For GPU runs, the best setup was:
 - **AmgX** (PCG solver, AMG preconditioner, Richardson smoother, direct solver for coarsest level) for the pressure equation
 - **Two-stage Gauss-Seidel** for Velocity and turbulence
- Mixed precision tested with native solvers only (no AmgX) **up to ~17%** performance improvement
- GAMG did not scale effectively from 24 to 48 GPUs
- Fastest run (48 GPUs) took **3615 seconds** total