

OpenFOAM Optimization Strategies for Next-Gen ARM CPUs

- Smart Platforms Lab, Huawei

- Anthony Berkow
- Igal Tsarfis
- Mark Wasserman
- Alon Zameret
- Yann Delorme
- Zhaohui Ding
- Weicheng Pei

- The 20th OpenFOAM Workshop, July 2025



OpenFOAM - Background

Open▽FOAM

❑ OpenFOAM in Huawei

- Huawei provides a full HPC stack
- OpenFOAM is extensively used by our customers

❑ OpenFOAM's growth relies on effectively leveraging massively parallel HPC architecture

❑ OpenFOAM's performance scaling is currently limited

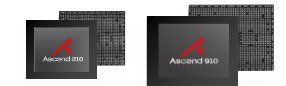
- Memory bandwidth: Irregular access patterns and low computational intensity → intra-node scalability
- Communication overhead: MPI communication dominates at large scale → inter-node scalability
- Computation efficiency: Complex data access patterns and lack of support for mixed-precision computation



CloudEngine (Switch)



OceanStor (Storage)



Ascend (AI)



IN200
(25/50/100GbE NIC)



Kunpeng 920 (CPU)

Investigated Optimization Goals

❑ Use modern architecture capabilities to reduce bottlenecks

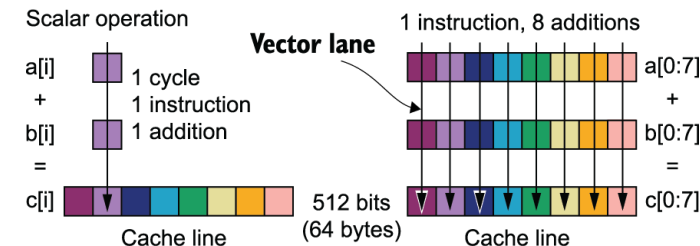
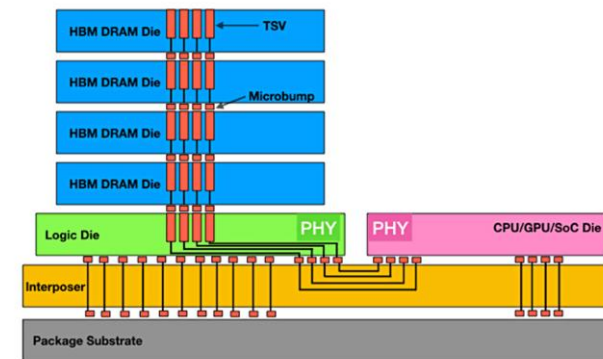
- Introduction of DDR5 and on-package memory → memory bound
- Many-core architecture → communication bound
- Vector/matrix SIMD units → computation bound

❑ Investigated directions

- Mixed-Precision (MxP) Solvers: Reduce memory bandwidth pressure and increase computational throughput
- Smart Memory Allocator (SMA): Efficient usage of all memory tiers (DDR/on-package memory) to accelerate memory access
- Vectorization: Usage of SIMD to accelerate computation
- Asynchronous Solvers: Reduce inter-process communication frequency

❑ This presentation will focus on the first two directions

HPL-MxP Benchmark Top 5 List, May 2024			
Rank (HPL-MxP)	Computer	Cores	Speedup of HPL-MxP over HPL
1	Aurora, intel/HPE Cray EX	8,159,232	10
2	Frontier, HPE Cray EX235a	8,699,904	8.3
3	LUMI, HPE Cray EX235a	2,752,704	6.2
4	Fugaku, Fujitsu A64FX	7,630,848	4.5
5	Leonardo, Bull Sequana XH2000	1,824,768	7.5



MxP: Overview

❑ Precision level is the number of significant digits a value accurately represents

- Low precision enables lower memory usage, faster computation and better hardware efficiency
- High precision ensures greater accuracy but incurs higher memory and computation overhead

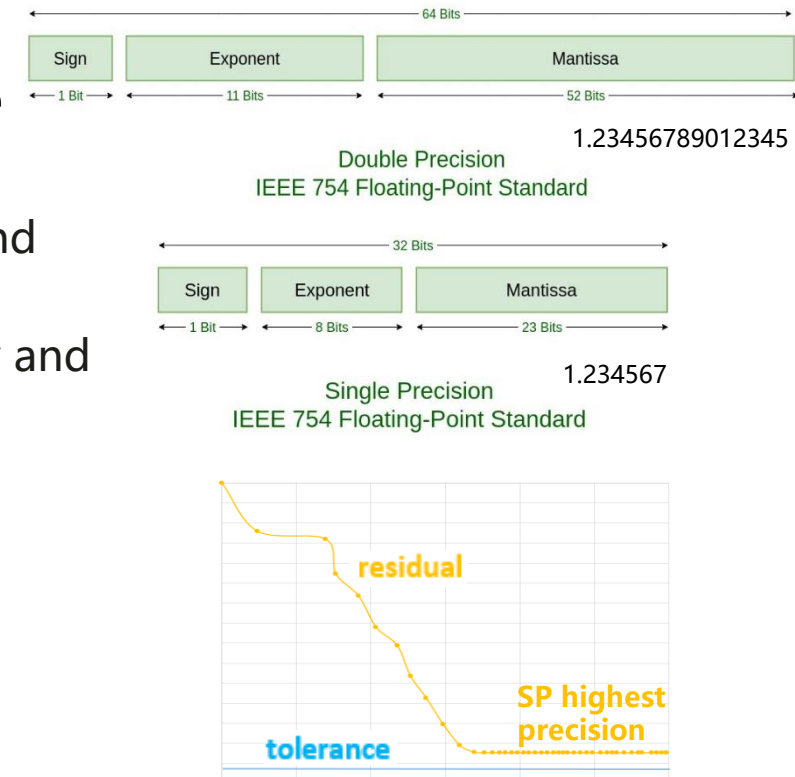
❑ MxP combines reduced precision and high precision

- Balance performance, memory usage and accuracy

❑ OpenFOAM supports the following precision levels

- Double-Precision (DP): Defines scalars as *double*
- Single-Precision (SP): Defines scalars as *float*
- Mixed-Precision (SPDP): Defines scalars as *float* but uses *double* for solvers
→ Compute in DP and truncate result to SP

❑ Our approach is opposite: Store in DP and solve in SP



MxP: Corrective Solver Algorithm

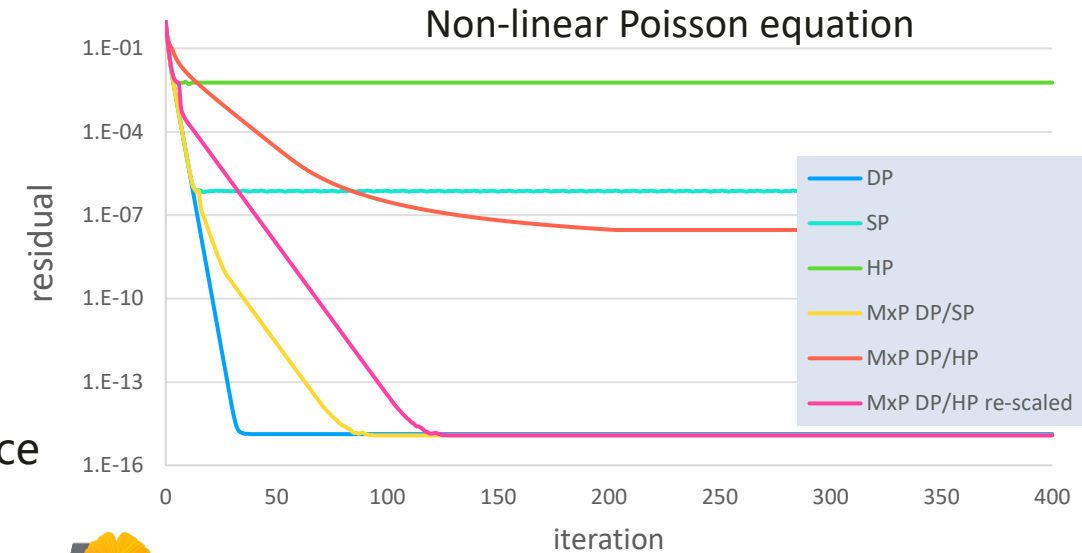
□ Linear solvers target $A \cdot x = b$ system of equations where

- $r = b - A \cdot \hat{x}$ is the residual for a solution \hat{x}
- Considered converged when $Norm(r) < tolerance$

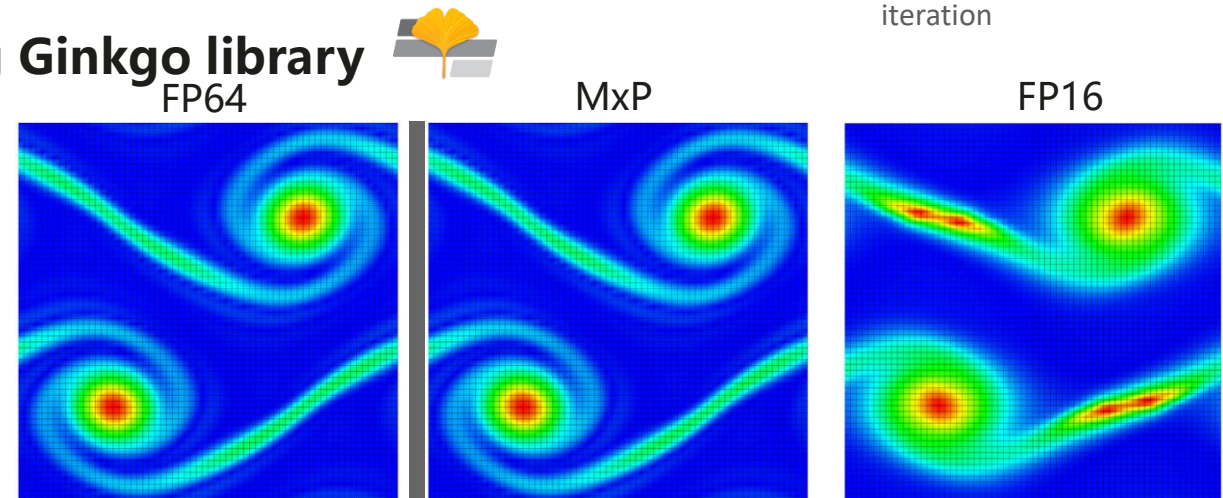
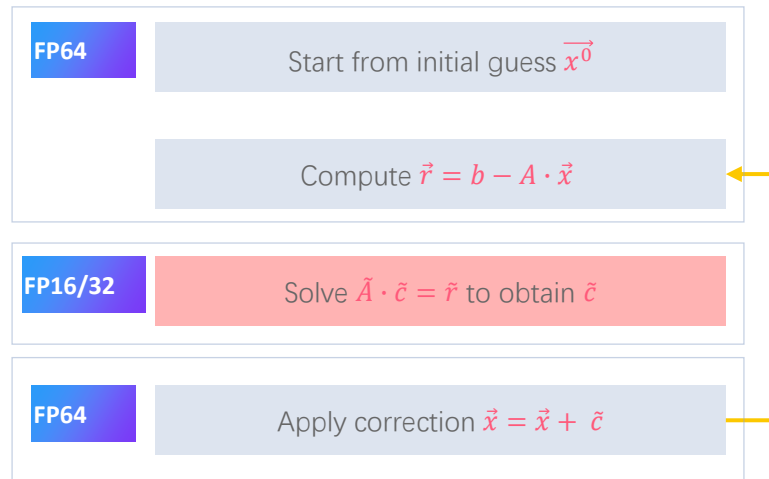
□ Developed MXPC solver that adds correction term

c such that $A \cdot (x + c) = b$

- We receive $A \cdot c = b - A \cdot \hat{x} = r$
- Obtains double-precision accuracy upon convergence



□ Preliminary PoC demonstrated using Ginkgo library



2D Double Shear Layer

MxP: Implementation in OpenFOAM

❑ Introduced mixed-precision (*MXPC*) solver

- Creates low-precision inner solver instance to enable MxP execution
- Templated numerous OpenFOAM classes to include support for solvers to have DP and SP instances
- Supports all segregated solvers, preconditioners and smoothers

```
template<class Type, class SType>
class TlduMatrix
```

```
typedef TGAMGSolver<floatScalar, floatScalar> GAMGSolverSP;
```

❑ Implemented as a standalone module

❑ Configured and used directly through the *fvSolution* dictionary

fields	scalarField, some of FieldFunctions
finiteVolume	coupledFvPatchField, processorFvPatchField, faceAreaPairGAMGAgglomeration
OpenFOAM/matrices	LUscalarMatrix, procLduInterface, procLduMatrix, LduInterfaceField, duInterfaceField, processorLduInterfaceField, lduMatrix, solverPerformance, scalarMatrices
OpenFOAM/matrices: Preconditioners	DIC, GAMG, none
OpenFOAM/matrices: Smoothers	GaussSeidel
OpenFOAM/matrices: Solvers	GAMG, PBiCGStab, PCG, diagonal, smoothsolver, MXPC
OpenFOAM/matrices: GAMG agglomerations	GAMGAgglomerateLduAddressing, GAMGAgglomeration, GAMGProcAgglomeration algebraicPair, dummy, pair, faceAreaPair, noneGAMGProcAgglomeration
OpenFOAM/matrices: GAMG interfaces	GAMGInterfaceField, processorGAMGInterfaceField, GAMGInterface
OpenFOAM/memory:	FieldFieldPrecisionAdaptor
OpenFOAM/primitives	Tensor

```
p
{
    solver      MXPC;
    tolerance   1e-7;
    relTol      0.01;

    SPsolver
    {
        solver      GAMG;
        smoother    GaussSeidel;
        tolerance   0;
        relTol      0.7;
        maxIter     50;
    }
}
```

MxP: 3D Cavity - MXPC validation (1 of 2)

❑ exaFOAM cavity3D benchmark (used for initial MXPC validation)

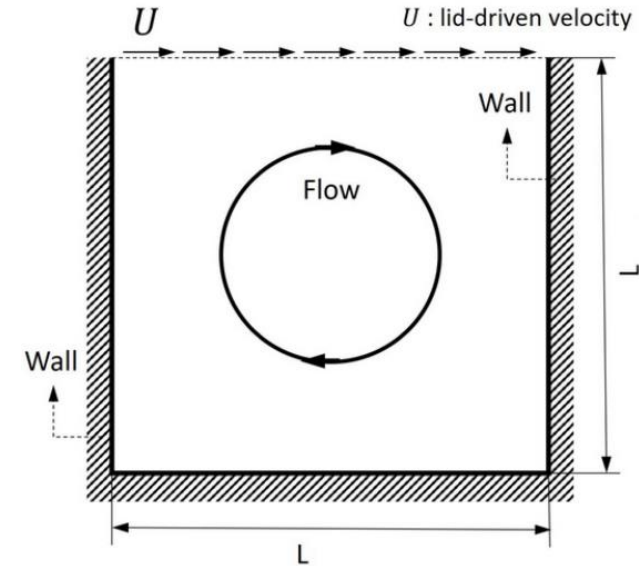
- Mesh size: 1M cells

❑ Setup:

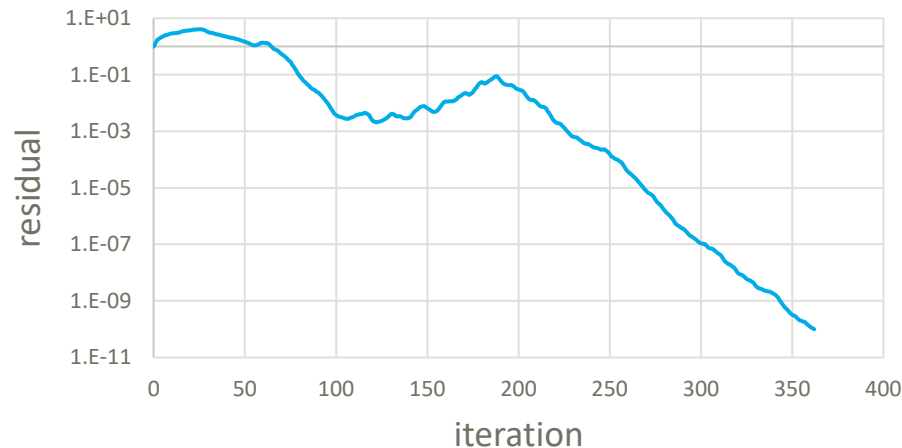
- Baseline: Pressure: DP PCG
- MxP: Pressure: MXPC + SP PCG

❑ Results:

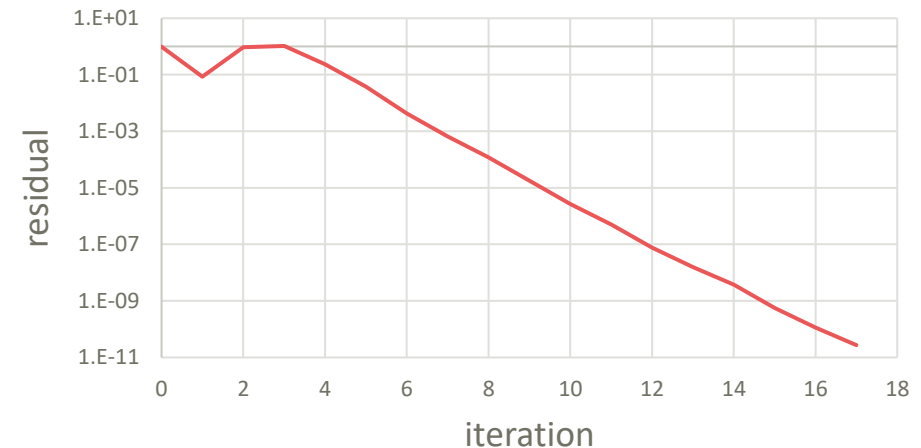
- MXPC retains the DP precision, no error accumulates over time steps.
- MXPC retains physical solution values (pressure and velocity).



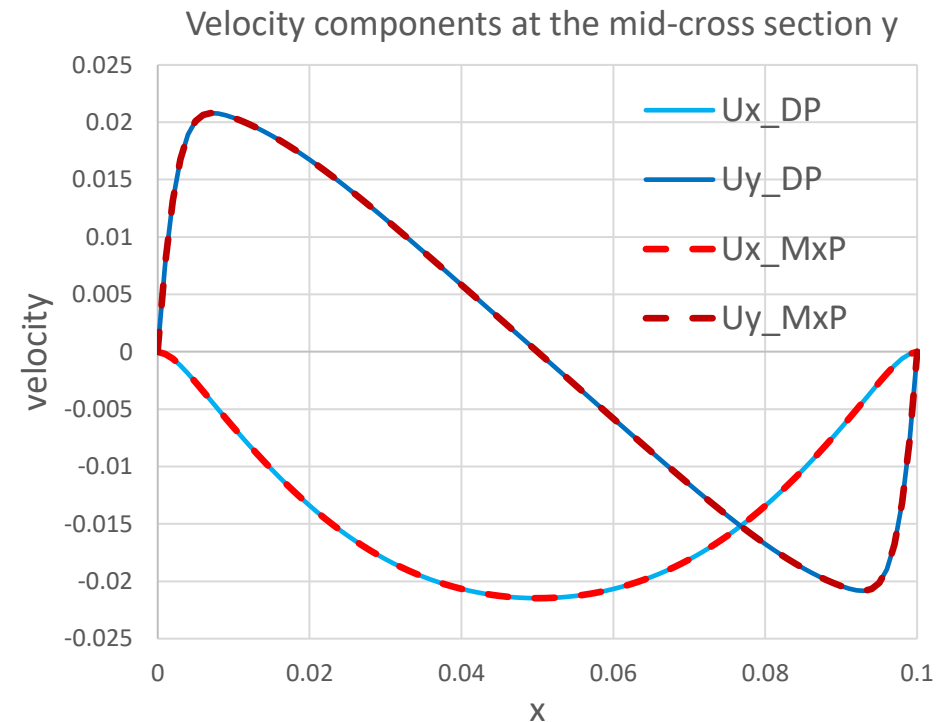
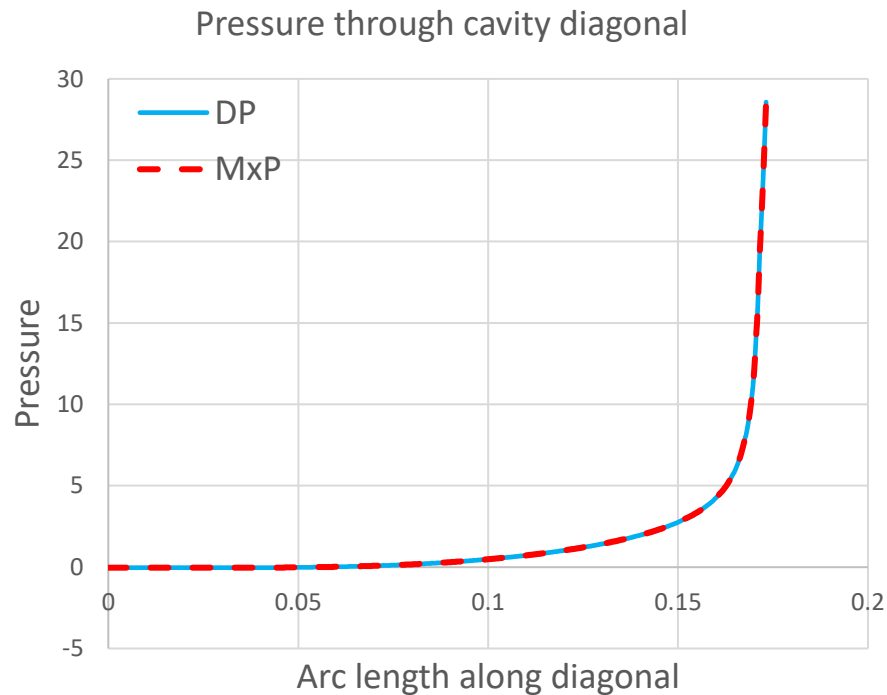
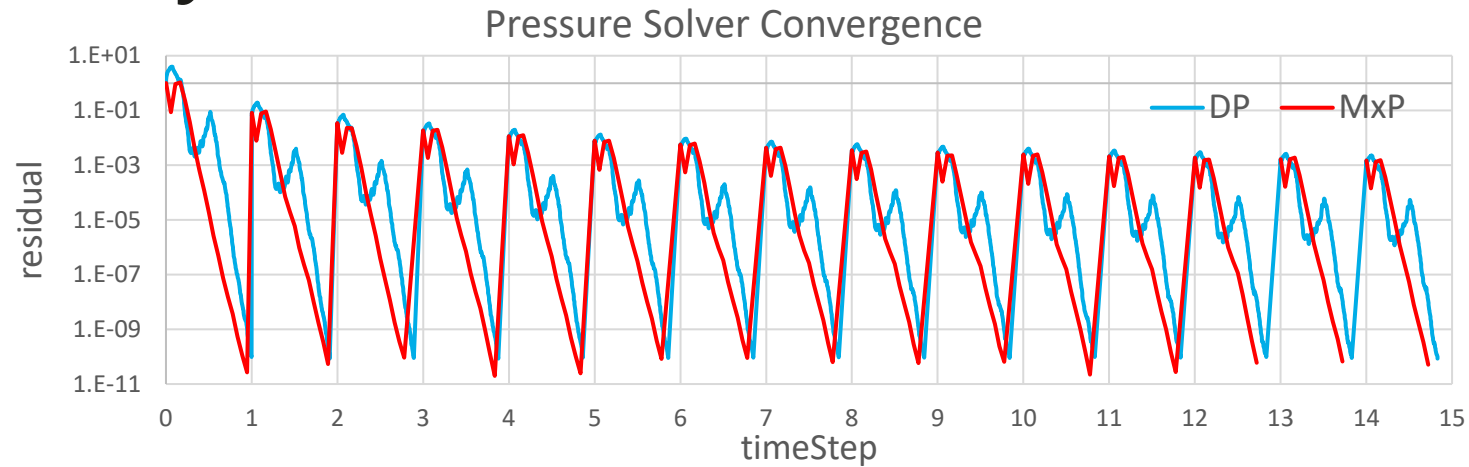
Baseline: convergence



MxP: convergence



MxP: 3D Cavity - MXPC validation (2 of 2)



MxP: DrivAer Benchmark with MXPC Solver

□ DrivAer B10 Benchmark

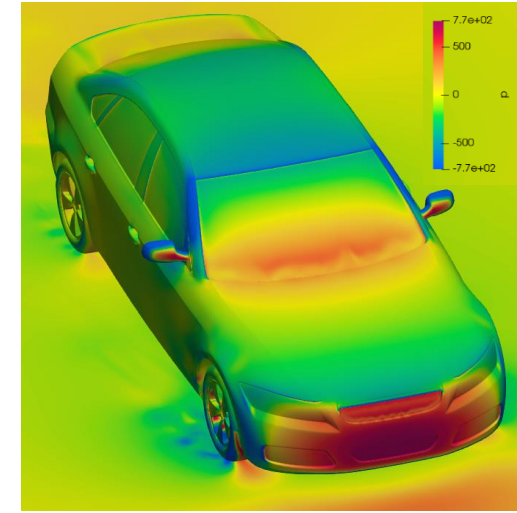
- DrivAer car model with stationary wheels and ground
- Solves **steady-state RANS** simulation with **simpleFOAM**
- Part of **OpenFOAM HPC Challenge 1** software track

□ Setup

- Baseline: Pressure: **DP GAMG** Velocity: **DP smoothsolver**
- MxP: Pressure: **MXPC + SP GAMG** Velocity: **MXPC + SP smoothsolver**

□ **MXPC obtained consistent speedup while maintaining DP accuracy**

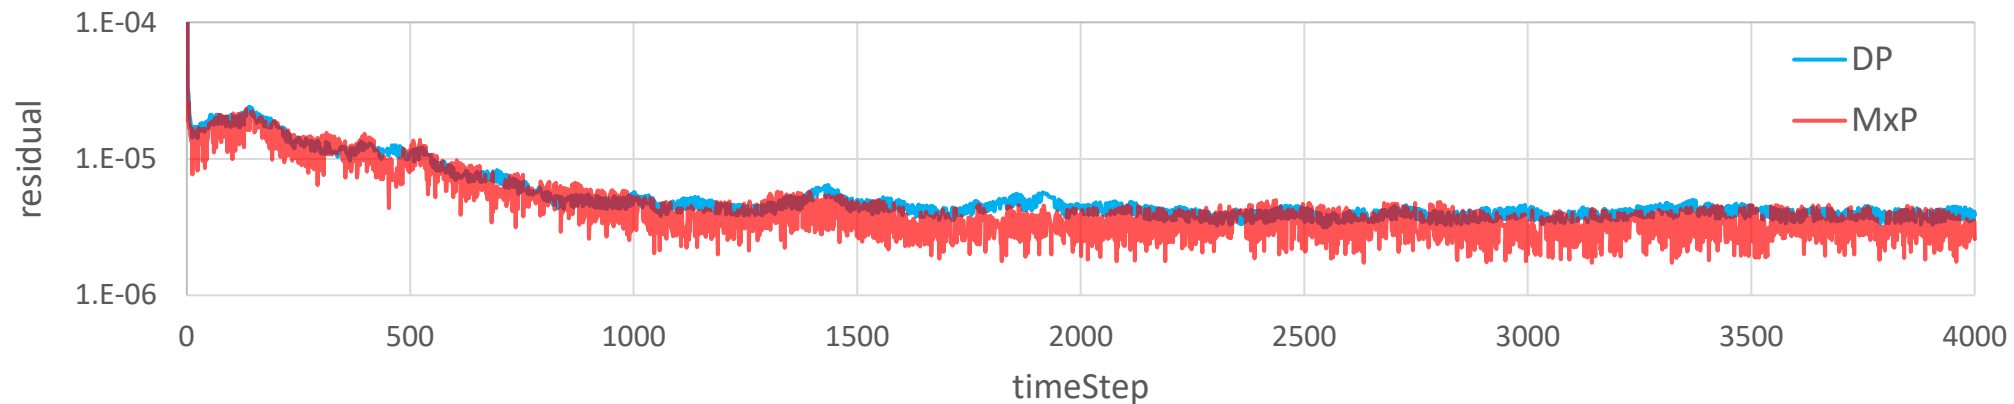
DrivAer mesh:	236M (fine)	p relTol:	0.01
Timesteps:	4000	U relTol:	0.1
Nodes:	8	Cores per node:	512
	DP	MxP	MxP Speedup
Total ClockTime (s):	22552	18098	1.25
Total solve time p:	7834	4013	1.95
Total solve time Ux+Uy+Uz:	3353	2643	1.27



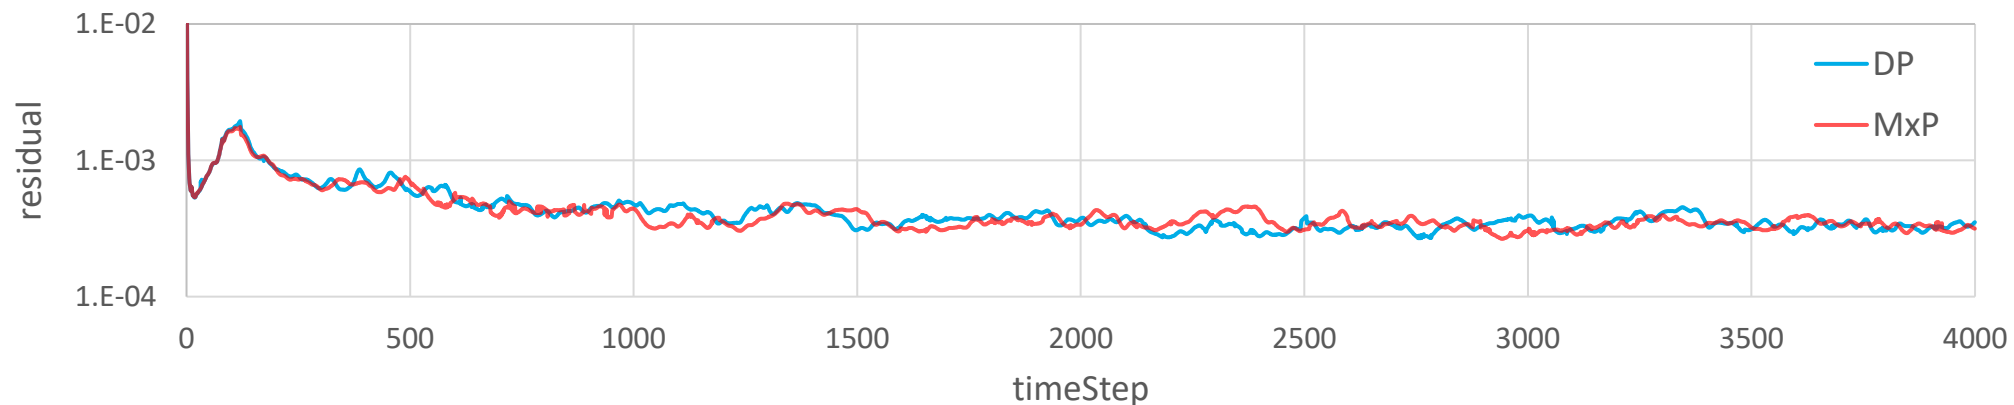
Hardware	
Node	2 LX2 CPU
LX2 pilot high-performance CPU	System on-chip >256 cores 2 Computing dies
Compute die memory	Off-die DDR On-package memory
Frequency	>1.3GHz
Interconnect	RDMA

MxP: DrivAer Benchmark - convergence

Convergence for p (relTol: 0.01)

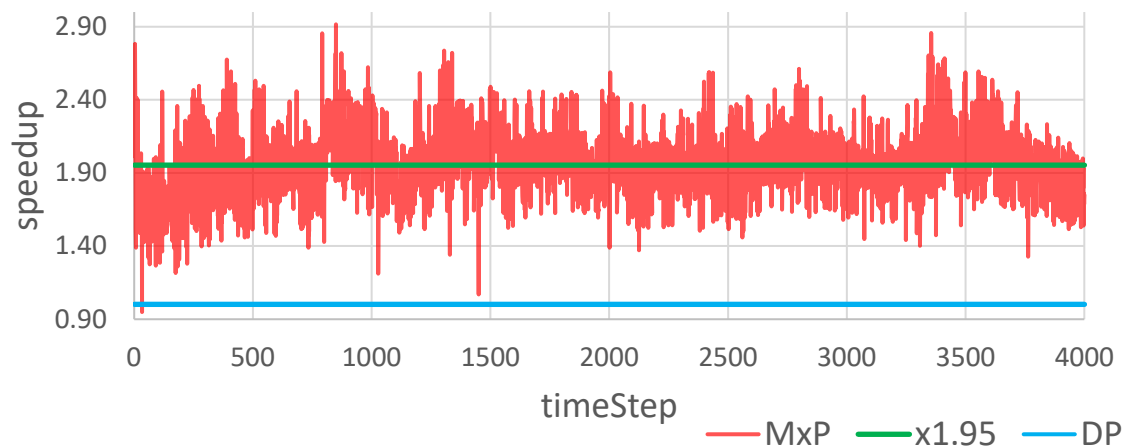


Convergence for U^* (relTol: 0.1)

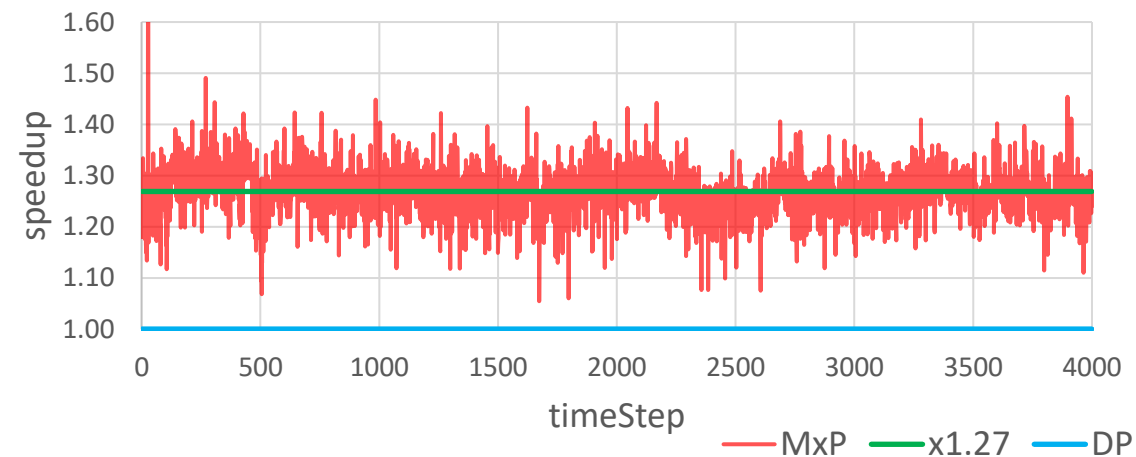


MxP: DrivAer Benchmark - MXPC speedup

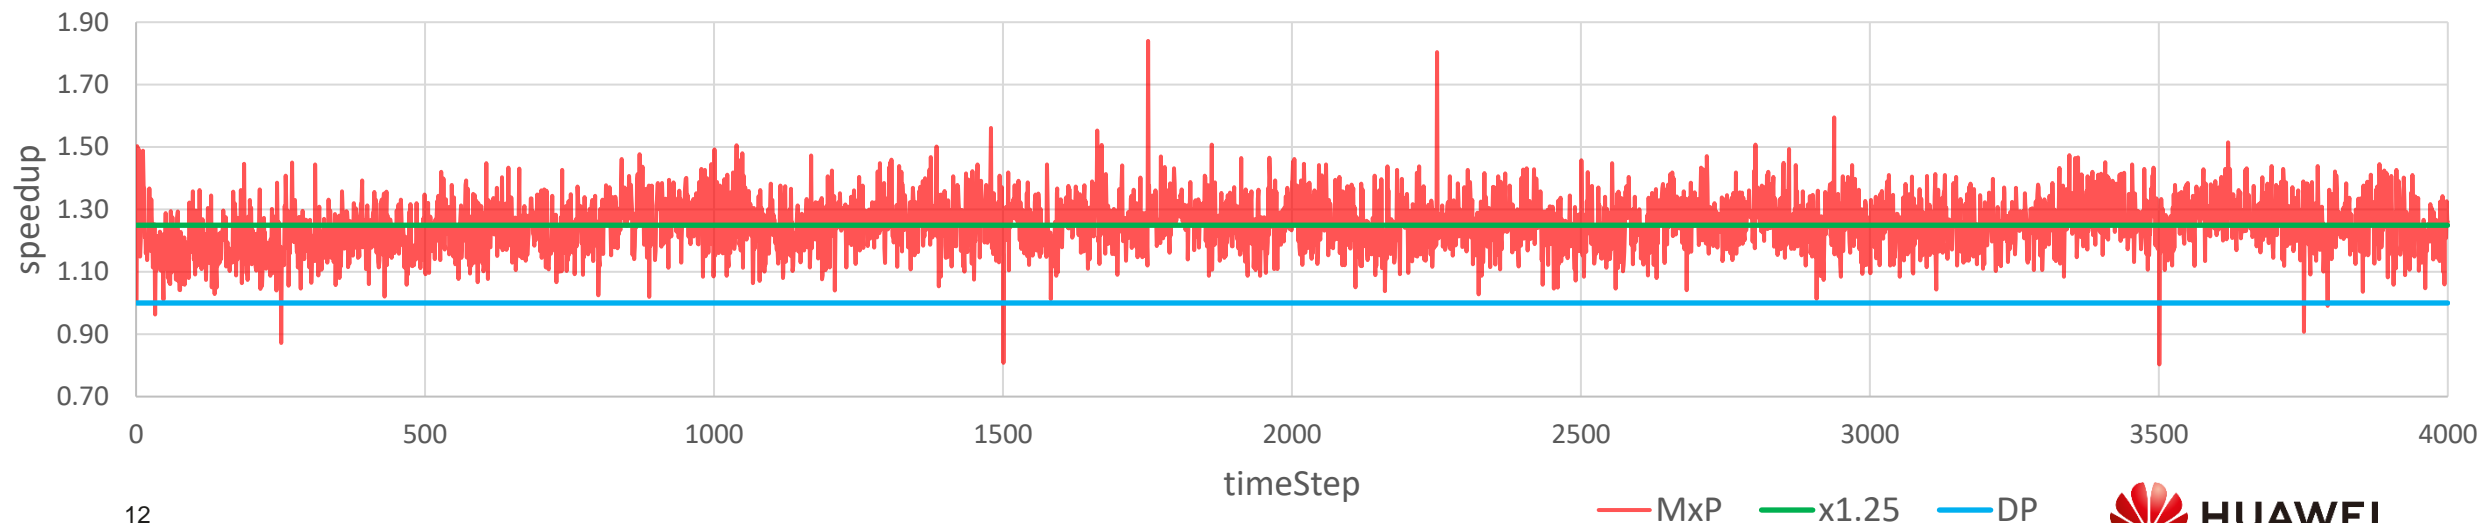
Speedup solving for p (relTol: 0.01)



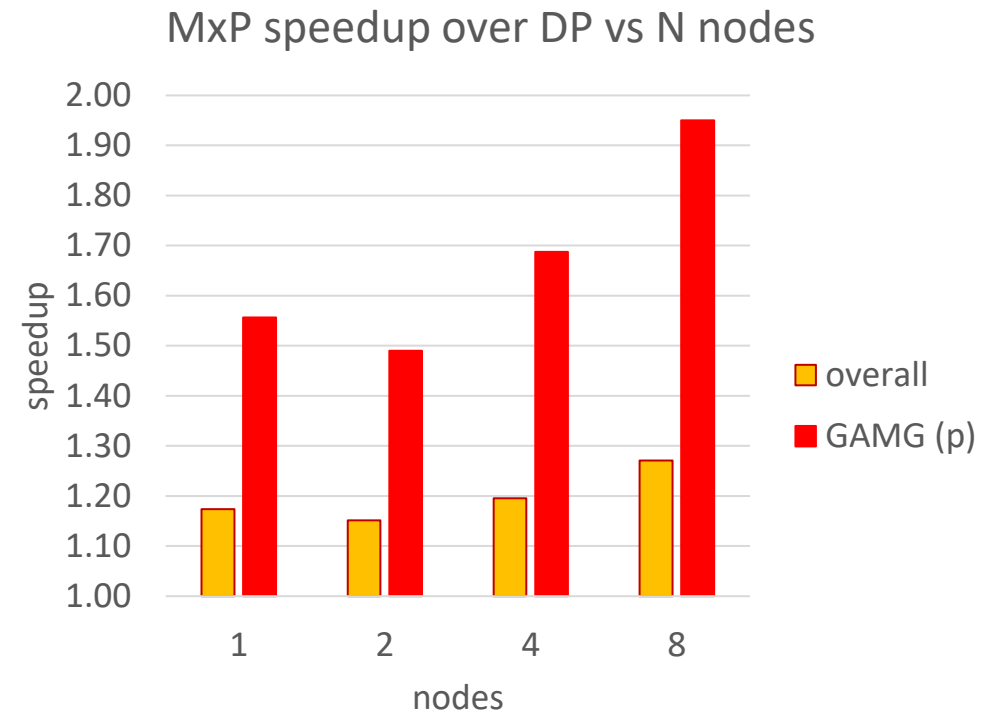
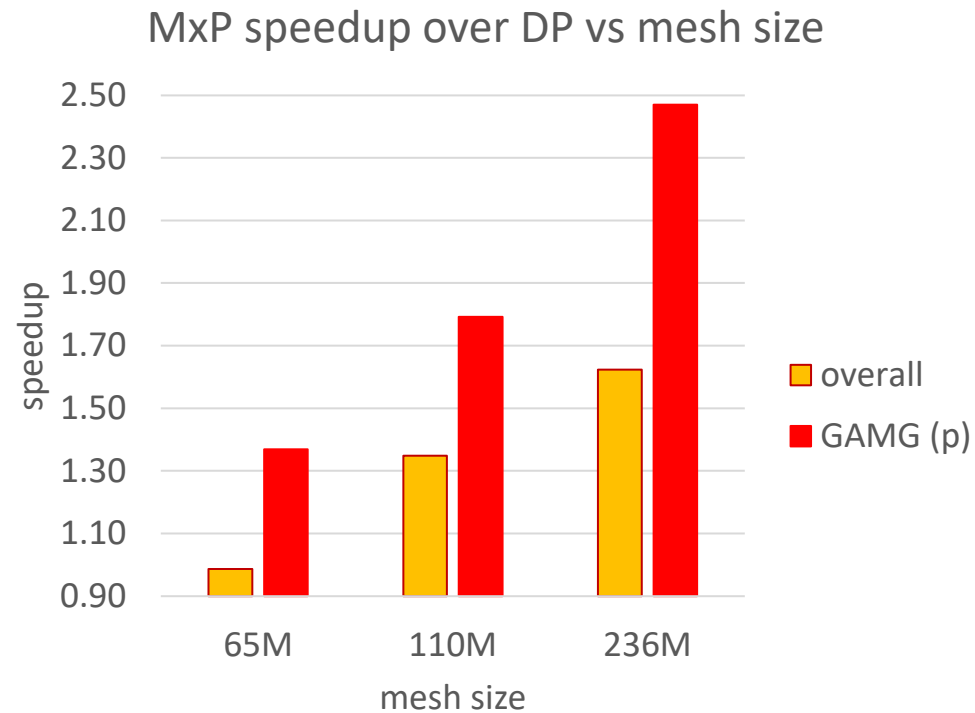
Speedup solving for U^* (relTol: 0.1)



Overall simpleFoam speedup



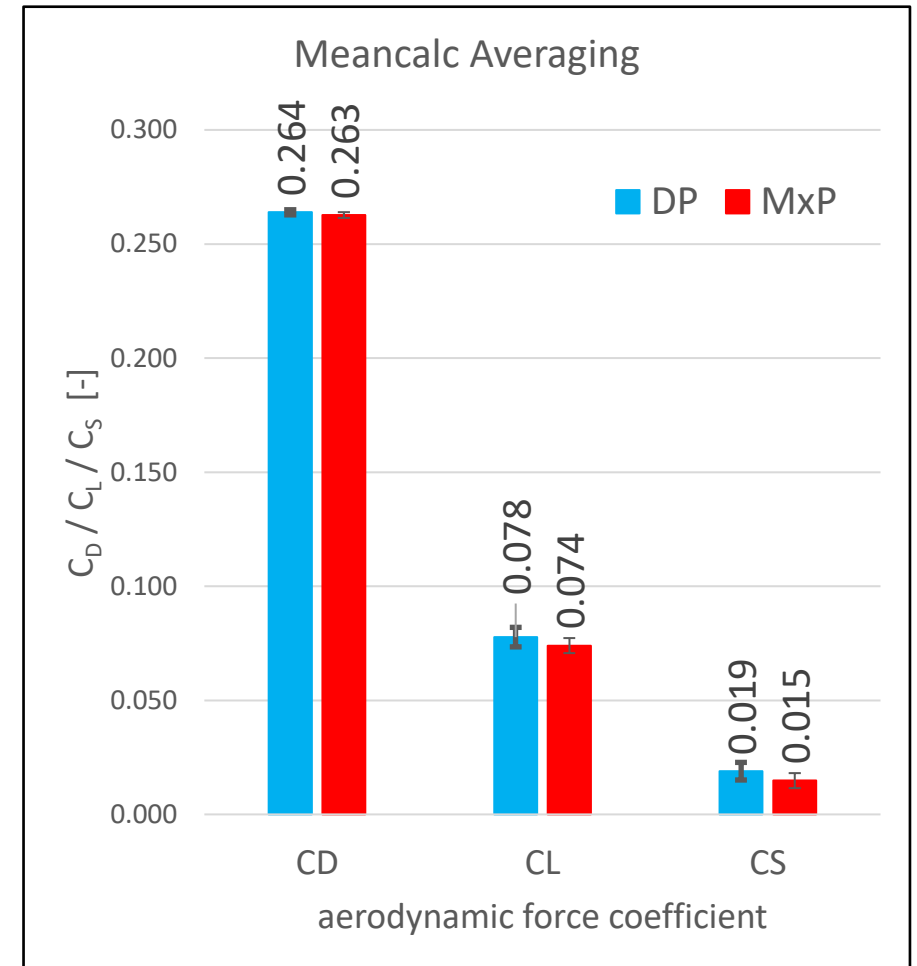
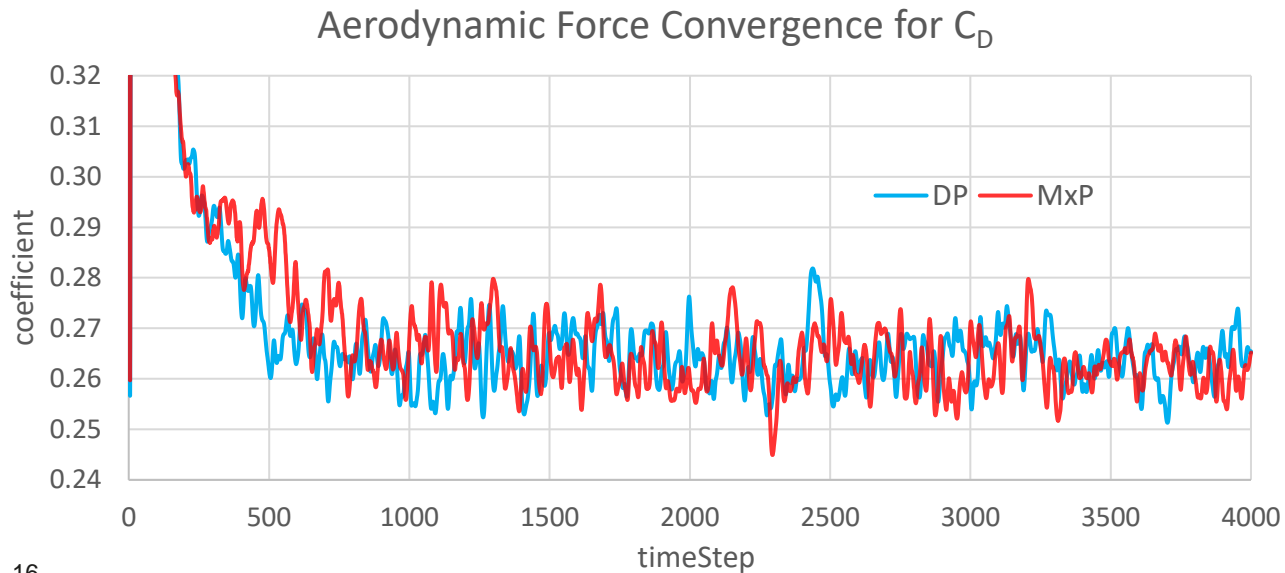
MxP: DrivAer Benchmark - MXPC speedup vs mesh size & N nodes



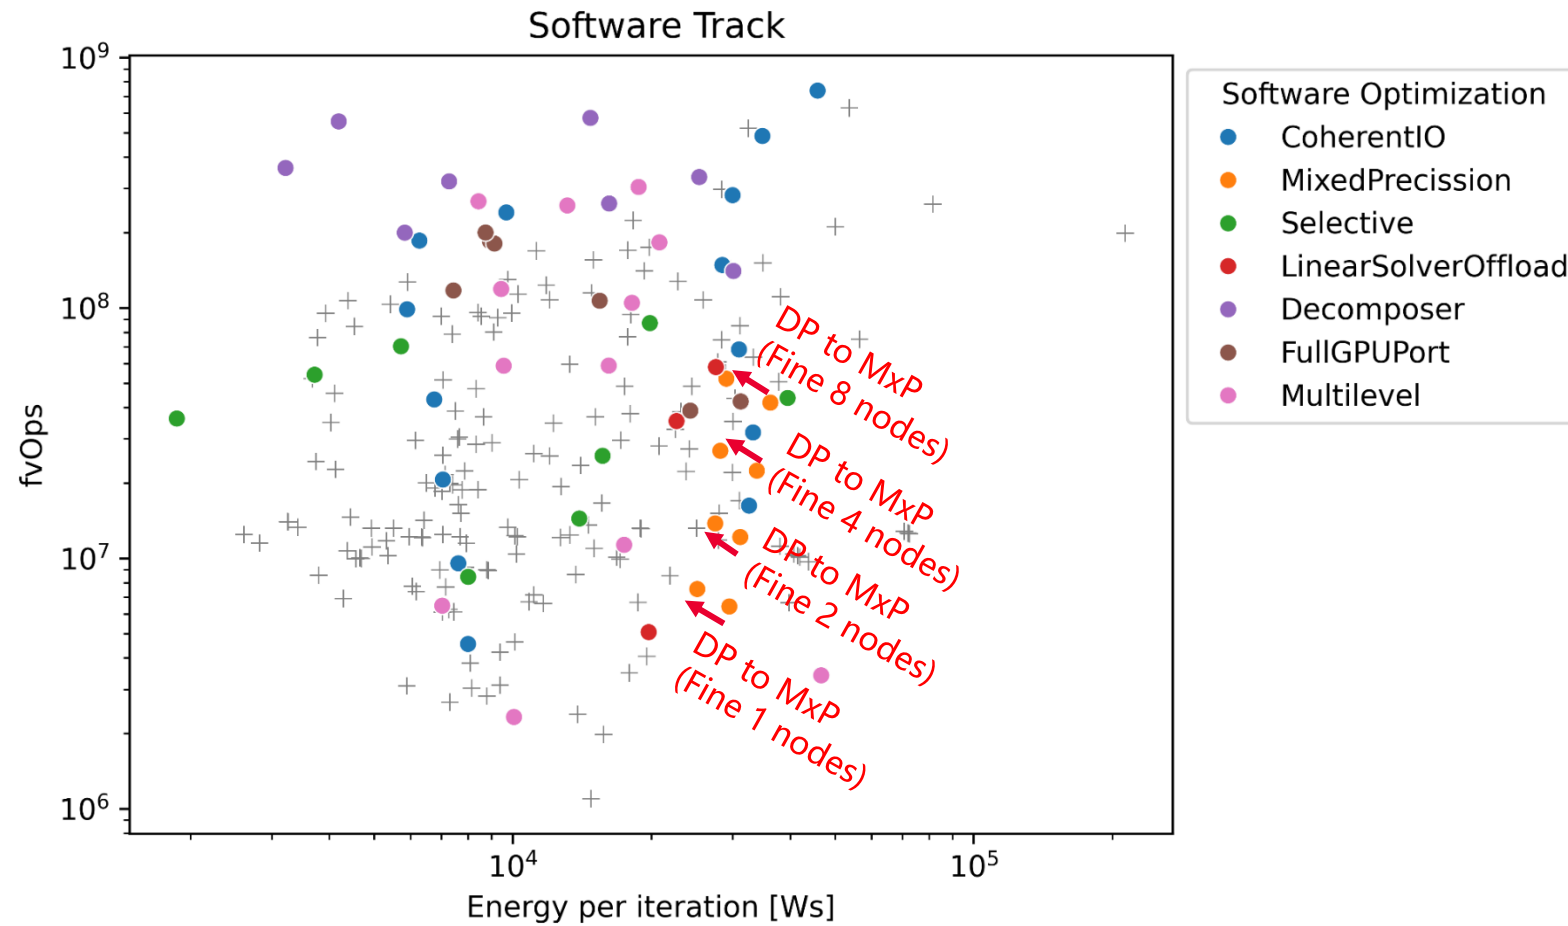
Note: Solving for p: relTol: 0.01 (inner relTol: 0.7)

MxP: DrivAer Benchmark - MXPC accuracy

Standardized Meancalc based evaluation of aerodynamic forces						
	C _D	DP C _L	C _S	C _D	MxP C _L	C _S
Total timesteps [nread]	4000	4000	4000	4000	4000	4000
Averaging start timestep [nskip]	1719	1719	1719	1723	1723	1723
Averaging samples [nused]	2281	2281	2281	2277	2277	2277
Mean value μ	0.2639	0.0778	0.0190	0.2627	0.0740	0.0149
Error Mean Value $\sigma(\mu)$	0.00055	0.00213	0.00194	0.00065	0.00165	0.00166
95% conf. int. on mean $2*\sigma(\mu)$	0.0011	0.0043	0.0039	0.0013	0.0033	0.0033
Standard deviation s	0.0050	0.0202	0.0087	0.0052	0.0180	0.0088
Error Standard Deviation $\sigma(s)$	0.0004	0.0015	0.0014	0.0005	0.0012	0.0012



MxP: DrivAer Benchmark - MXPC results



Smart Memory Allocator (SMA) - Overview

❑ OpenFOAM performance is often memory-bound due to

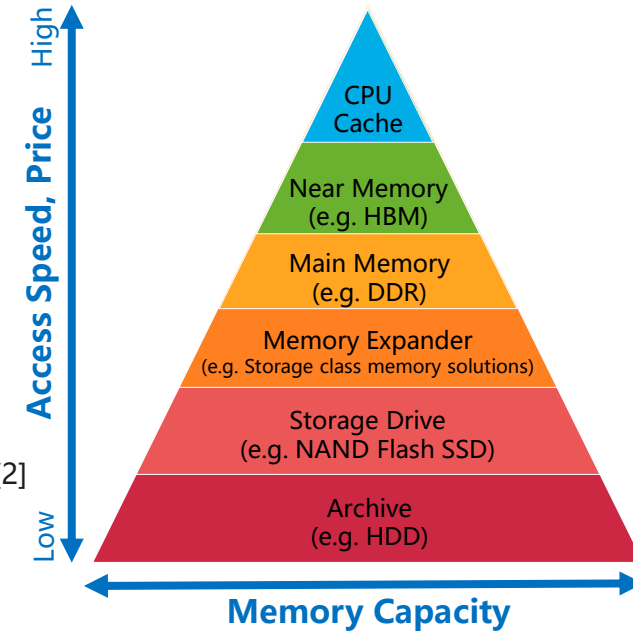
- Irregular access pattern to large field arrays
- Low computational intensity and lack of kernel fusion

❑ Leveraging faster memory tiers can alleviate this bottleneck

- Allow data-delivery rates to better align with computational throughput
- Demonstrated by AMD^[1] on MI300A while using only 128 GB unified HBM^[2]
- Larger mesh → more machines → higher cost and energy consumption

❑ Smart Memory Allocator (SMA)

- Control memory tier assignment for each allocated buffer
- Memory bandwidth-bound kernel buffers → higher memory tier
- Maximize performance gain with limited capacity of higher-tier memory
- Obtain comparable performances while minimizing required number of machines
- Fewer machines required → reduce communication overhead



Smart Memory Allocator (SMA) - Implementation

❑ Based on *feature-memory_pool* branch (Mark Olesen)

- Memory resources (Lists/Matrix derived types) are management by **Umpire** library
- Utilize **pre-allocated memory buffer** for constructed variables using **single Umpire allocator**
- Acts as **abstraction layer** between application and underlying memory tier



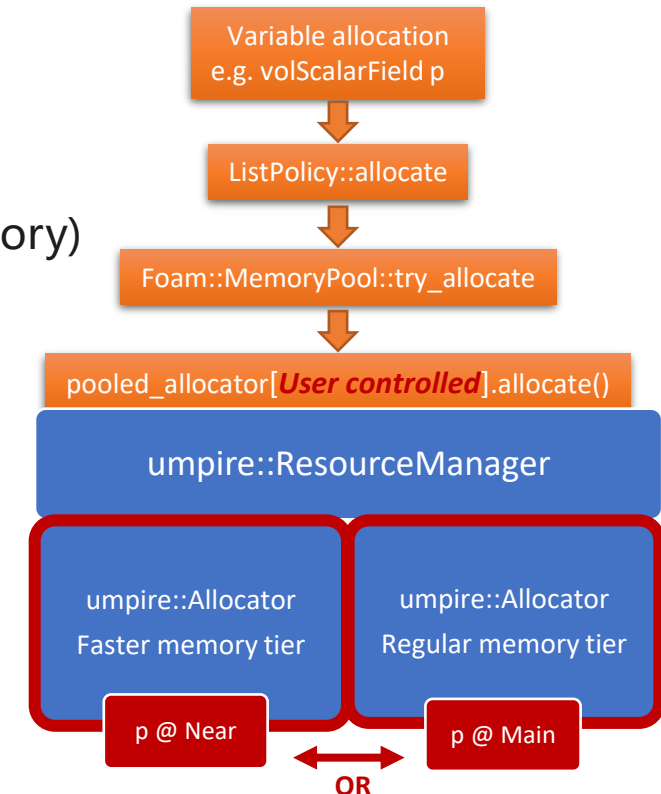
❑ Extended framework to support **two Umpire allocators**

- Utilize Umpire **built-in NUMA** strategy
- Each allocator **manages a different memory-pool tier** (Near / Main memory)

❑ Variable allocation is directed to desired memory pool based on:

- Current **compute kernel** being executed
- **User-defined assignment** of each **compute kernel** using environmental variables

❑ Choose the best **cost effective allocation** for given **mesh size** and **available resources**



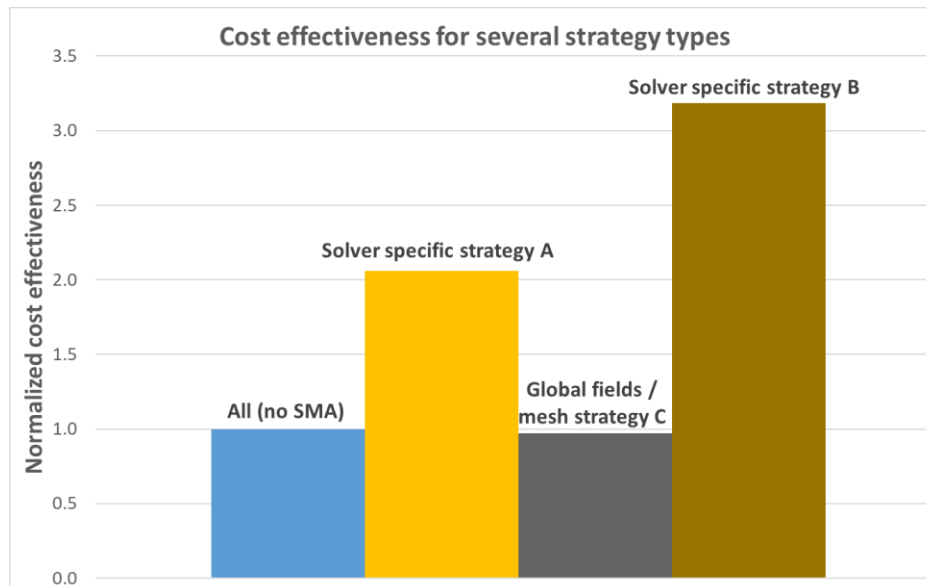
Smart Memory Allocator (SMA) - Results

❑ Test Case

- DrivAer B10 benchmark used as part of the OHC-1 evaluation (coarse/medium/fine mesh)
- Executed on LX2 pilot high-performance CPUs
- Umpire allocators assigned to DDR and on-package memory tiers

❑ Cost-effectiveness metric

- Decide which simpleFoam variables will be allocated on DDR / on-package memory tier
- Aim to effectively utilize on-package memory capacity to minimize run time
- Compared multiple strategies with varying memory demands (All > B > A > C)



$$\text{Cost effectiveness} = \frac{\text{Acceleration}}{\text{Ratio of used on-package memory}}$$

$$\text{Acceleration} = \frac{\text{Reference wall clock time (all variables on DDR)} - \text{Wall clock time for given strategy}}{\text{Reference wall clock time (all variables on DDR)}}$$

$$\text{Ratio of used on-package memory} = \frac{\text{Variables allocated on on-package memory for given strategy [GB]}}{\text{All variables allocated on on-package memory [GB]}}$$

$$\text{Normalized cost effectiveness} = \frac{\text{Cost effectiveness for given strategy}}{\text{Cost effectiveness of reference (all variables on on-package memory)}}$$

Smart Memory Allocator (SMA) - Results

❑Speedup

- Consistent **speedup over DDR**
- **Retains** performance compared to pure on-package memory execution

❑Enabler

- Major **enabler** for scarce resources
- Coarse mesh on a single machine, Medium and fine meshes factors required #machines
- Utilization of **on-package memory when capacity is not sufficient for pure run**

❑Energy efficient

- **Fewer machines required + Retaining performance → High energy efficiency**

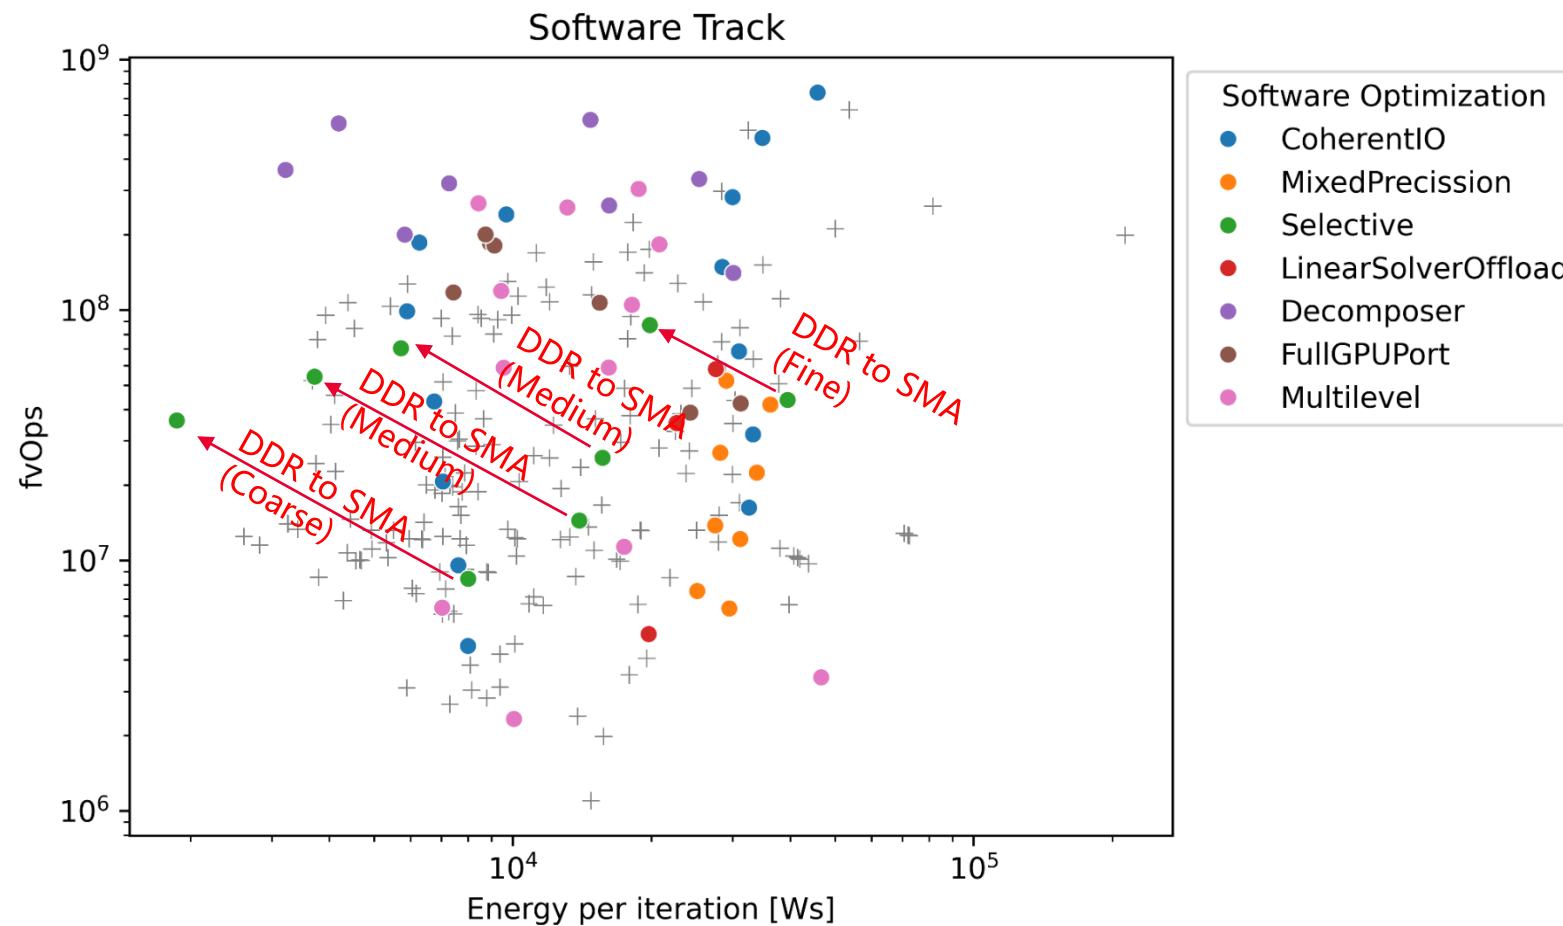
				Total consumed energy [kW*s/step]	
				Smart Memory Allocator	pure on-package memory
Mesh size	#Nodes	#Cores	WCT Speedup SMA vs. pure DDR		
Coarse (~65M)	1	>512	4.3X	1.9 (1 node)	3.7 (4 nodes)
	1	>512	1.3X	11.5 (1 node)	6.6 (8 nodes)
Medium (~110M)	2	1024	3.7X	3.7 (2 nodes)	6.6 (8 nodes)
	4	2048	2.7X	5.7 (4 nodes)	6.6 (8 nodes)
Fine (~236M)	8	4096	2X	19.8 (8 nodes)*	17.7 (14 nodes)

23



*Fine mesh energy demonstrates **further scaling** can be made to allow **higher efficiency** (8->10 nodes), exploiting the still existing memory bandwidth bottleneck

Smart Memory Allocator (SMA) - Results



Future Directions

❑ Integrate selective memory allocation within MxP solver

- Allocate only single-precision buffers on faster memory tier
- Accelerate SP solver and support larger simulations on limited on-package memory capacity

❑ Integrate asynchronous solvers within OpenFOAM

- Reduce inter-process communication frequency → improve scalability
- Obtain convergence in less iterations → reduce computation

❑ Evaluate options for contributing the presented optimizations to the OpenFOAM community

