

# Software Track Submission

Gregor Olenik


- Our focus is offloaded linear solver
- No commonly accepted GPU OF version available that allows to offload matrix assembly
- Plugin based approaches face Over- and Undersubscription Challenge [1]
- Investigated individual partitioning for CPU and GPU workloads

- For efficient repartitioning we use multilevel decomposition, typical  $\{n\_nodes, n\_gpu, n\_cores\}$  to avoid “holes” in partition
- Automated pre-processing procedure OBR [2], also used for baseline CPU runs typical  $\{n\_nodes, n\_sockets, n\_cores\}$  -> Software Track

4	5	5	6	7	8
1	2	3	2	3	4
0	3	4	5	0	1
6	7	1	2	6	7
4	5	0	3	4	5
0	1	3	0	1	2


4	5	5	6	7	8
1	2	3	2	3	4
0	3	4	5	0	1
6	7	1	2	6	7
4	5	0	3	4	5
0	1	3	0	1	2

- Changing partitioning from hierarchical to multilevel alone improved performance by approximately 40%
- This be potentially improved further by (at least for GPUs) by considering mapping of subdomains to devices to maximize bandwidth




# Results Offloading GPU CG solver


- OpenMPI is susceptible to performance degradation when oversubscribing (more than 1 MPI rank per Device) GPUs



# CG vs Multigrid

- Multigrid (GAMG) outperforms any CG implementation (CPU/GPU)
- There might be a range between 16 - 32 nodes, where CG is faster. But in general GPUs profit from more DoF per device.
- Thus the obvious choice would be to use a GPU Multigrid solver/





- Ginkgo now features different flavours of (algebraic) Multigrid
  - Solver, preconditioner, distributed, local as RAS
- However, the multigrid parameter space is large. CPU based setting wont necessarily work well for GPUs (fewer levels)
- Reuse of geometric restriction and prolongation